Simulation for a priority queue with N events

Chenwei Yue
GT ID 903243179

The definitions of most struct and functions are similar with that in the simulation program with three stations. But in this program, the events only have two attributes: a pointer to the next event and a timestamp. The timestamp of events are generated by the urand() function. After adding N events into the priority queue, the timer starts. The timer stops after 5000 loops of event insert and removal. I created two variables: struct timeval start and struct timeval finish. These variables have two important attributes: tv_usec and tv_sec, which records the microsecond and second of system time, respectively. So the operation time for this 5000 loops is:

duration=(finish.tv_sec -start.tv_sec) *1000000+(finish.tv_usec-start.tv_usec)

Here I used gettimeofday() function to record the system time.

When N increases, the average operation time for one iteration also increases. The results are shown below. (The unit of time is microsecond)

N:	total time:	average time:
1000	4003.000000	0.800440
3000	34524.000000	6.903419
5000	81558.000000	16.308338
7000	109580.000000	21.911618
9000	157612.000000	31.516097
11000	220656.000000	44.122376
13000	327230.000000	65.432913
15000	436309.000000	87.244351
17000	561897.000000	112.356929
19000	618559.000000	123.687063
21000	692490.000000	138.470306
23000	832088.000000	166.384323
25000	923653.000000	184.693661
27000	1085269.000000	217.010398
29000	1166826.000000	233.318536
31000	1348454.000000	269.636873
33000	1694697.000000	338.871626
35000	1711710.000000	342.273545
37000	1751237.000000	350.177365
39000	2782468.000000	556.382324

The graph for (N,average operation time) is:

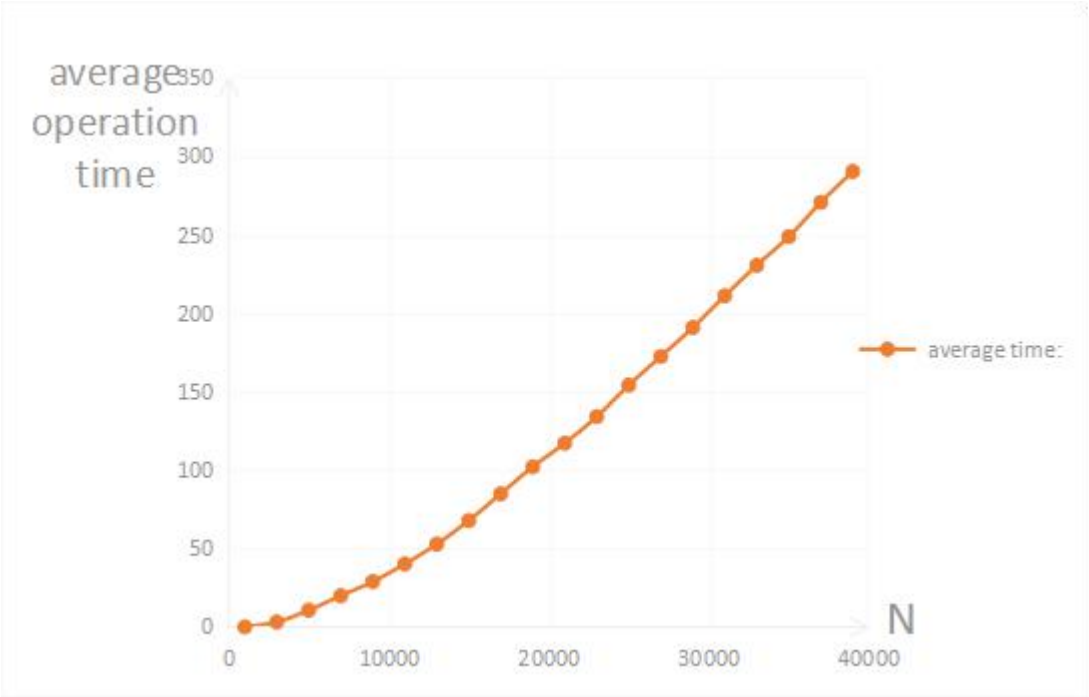

According to this graph, we can conclude that the average operation time is approximately linearly related to N.